

Advanced Django Architecture

Rami Sayar - @ramisayar

Technical Evangelist

Microsoft Canada

~~Advanced~~ Scalable Django Architecture

Rami Sayar - @ramisayar

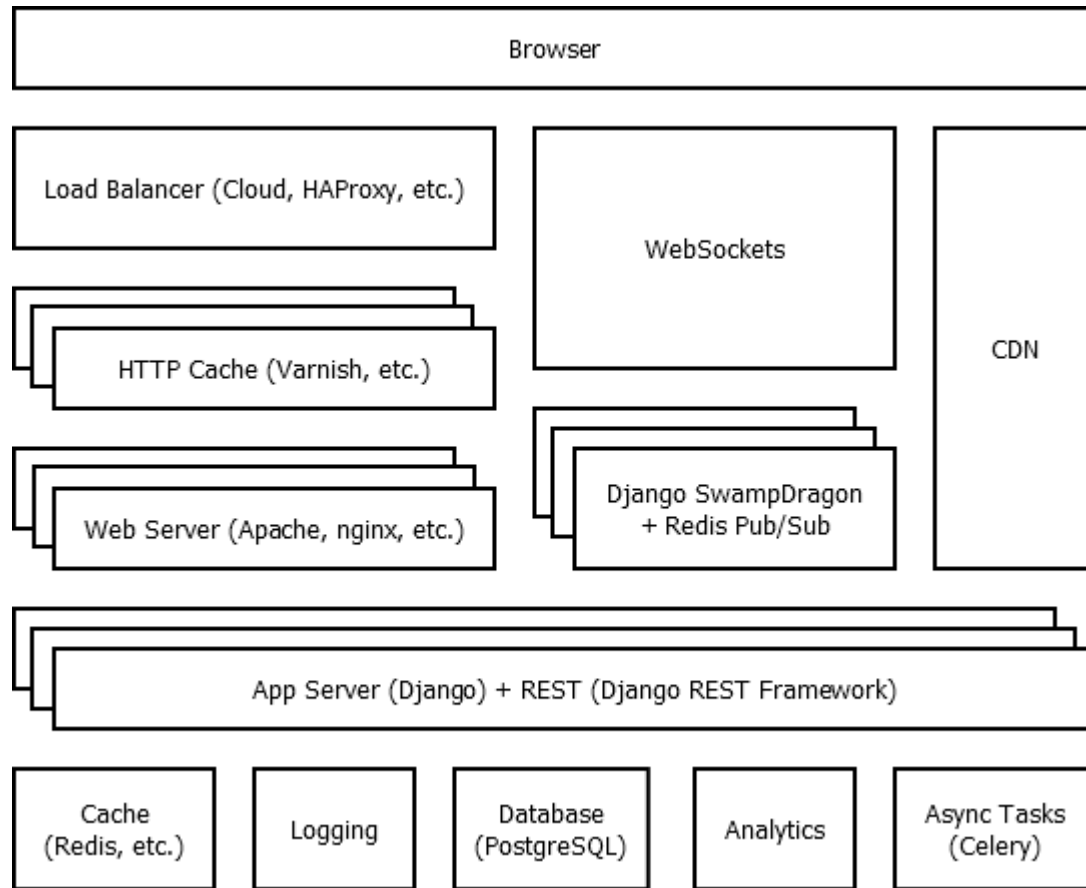
Technical Evangelist

Microsoft Canada

Agenda

- Advanced Django Techniques & Patterns
- Structuring Complex Applications
- Scaling Django Applications
- Django Applications in Production

What is a production web app?



- Django is surrounded by tons of applications, services and code in production.
- Advanced Django architecture is needed to support Django in production without failing hard.

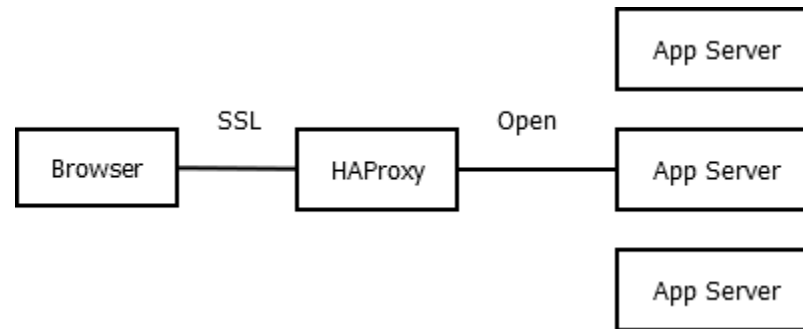
Load Balancing Django

Load Balancing

- Load balancing Django improves the performance and reliability by distributing traffic.
- Becomes an interesting problem when you have to consider:
 - SSL Termination
 - Sticky Sessions?
 - WebSockets?

Load Balancing = Django + HAProxy + SSL

- HAProxy gained native SSL support in HAProxy 1.5.x => June 2014!
- SSL Termination:

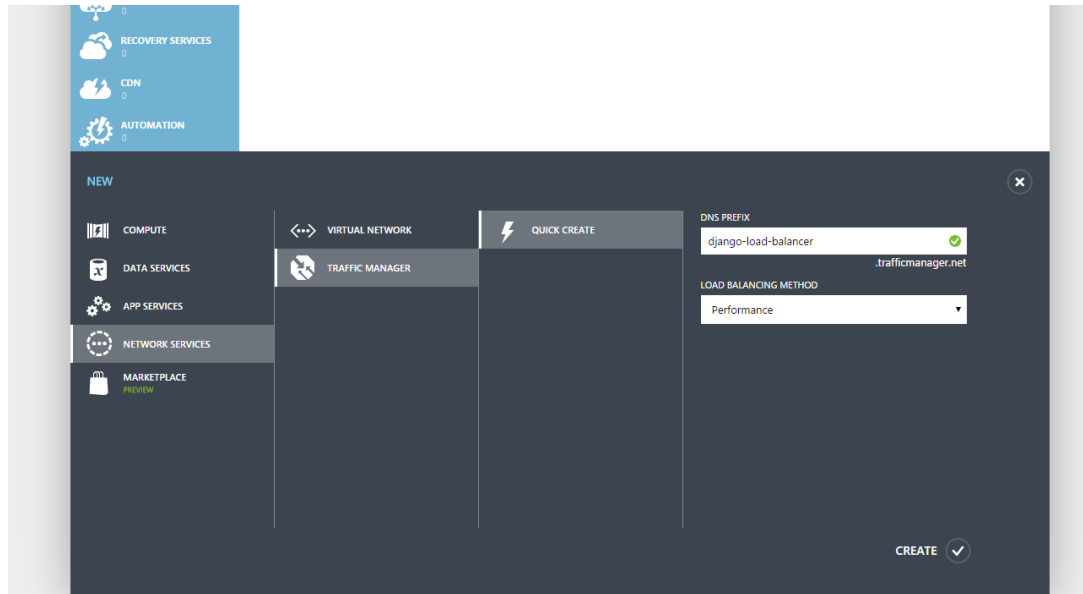


- Fairly straightforward to set up HAProxy.
 - Read: [How To Implement SSL Termination With HAProxy](#)

Load Balancing Proxying & Remote Addr

```
class SetRemoteAddrFromForwardedFor(object):  
    def process_request(self, request):  
        try:  
            real_ip = request.META['HTTP_X_FORWARDED_FOR']  
        except KeyError:  
            pass  
        else:  
            # HTTP_X_FORWARDED_FOR can be a comma-separated list of IPs.  
            # Take just the first one.  
            real_ip = real_ip.split(",")[0]  
            request.META['REMOTE_ADDR'] = real_ip
```


Azure Traffic Manager



- High Performance Load Balancer + DNS Support.
- Sticky Sessions don't exist.
- HTTPS Supported.

Caching Django

Caching = Django + Varnish

- Varnish has a tendency to just cache EVERYTHING which can be frustrating for the User, unless you set it up properly.
- Varnish does not cache content & cookies.
- Varnish and Django CSRF protection gotcha.
- Varnish does not deal with HTTPS (hence the need of HAProxy in front of Varnish).

Caching = Django + Varnish

- Django-Varnish – “It allows you to monitor certain models and when they are updated, Django Varnish will purge the model's `absolute_url` on your frontend(s). This ensures that object detail pages are served blazingly fast and are always up to date.”
- <https://github.com/justquick/django-varnish>
- Project Idea: Deeper Django/Varnish integration.

Caching = Django + Varnish

- <https://www.varnish-cache.org/>
- <http://www.isvarnishworking.com/>
- <http://yml-blog.blogspot.ca/2010/01/esi-using-varnish-and-django.html>
- <http://chase-seibert.github.io/blog/2011/09/23/varnish-caching-for-unauthenticated-django-views.html>
- <http://blog.bigdinosaur.org/adventures-in-varnish/>
- <http://www.nedproductions.biz/wiki/a-perfected-varnish-reverse-caching-proxy-vcl-script>

Serving Django

Web Server = Django + Apache

- Forced to use Apache? Your options:
 - mod_python (No)
 - mod_wsgi (Yes)
 - FastCGI (Deprecated since 1.7)
 - Phusion Passenger (Yes)

Web Server = Django + Apache

- mod_wsgi
 - Decent performance on the Apache side.
 - You can setup Apache any way you like.
 - You can setup mod_wsgi to use virtualenv.

Web Server = Django + Apache

```
import os, sys, site
```

```
# Add the site-packages of the chosen virtualenv to work with  
site.addsitedir('~/.virtualenvs/myprojectenv/local/lib/python2.7/site-packages')
```

```
# Add the app's directory to the PYTHONPATH  
sys.path.append('/home/django_projects/MyProject')  
sys.path.append('/home/django_projects/MyProject/myproject')
```

```
os.environ['DJANGO_SETTINGS_MODULE'] = 'myproject.settings'
```

```
# Activate your virtual env  
activate_env=os.path.expanduser("~/virtualenvs/myprojectenv/bin/activate_this.py")  
execfile(activate_env, dict(__file__=activate_env))
```

```
import django.core.handlers.wsgi  
application = django.core.handlers.wsgi.WSGIHandler()
```

Web Server = Django + Alternatives (aka. Nginx)

- Nginx + uWSGI

[uwsgi]

touch-reload = /tmp/newproject

socket = 127.0.0.1:3031

workers = 2

chdir = /srv/newproject

env = DJANGO_SETTINGS_MODULE=newproject.settings

module = django.core.handlers.wsgi:WSGIHandler()

Web Server = Django + Alternatives (aka. Gunicorn)

- Nginx & Gunicorn
 - In theory, you should use Nginx as a reverse proxy here, to serve static files if you're not uploading everything to a CDN or a separate media server. If you are, you can use HAProxy and Gunicorn as a standalone webserver, one less dependency to care about.
 - Install Gunicorn directly inside your virtual environment.
 - Use Gaffer to monitor the Gunicorn.

Web Server - References

- <http://blog.kgriffs.com/2012/12/18/uwsgi-vs-gunicorn-vs-node-benchmarks.html>

REST APIs & Django

Use Django REST Framework!

Django REST Framework vs Django Tasty Pie

Cache & Django

Caching in Production

- Caching – Django has multiple options:
 - **Redis.** Make sure you use Hiredis and django-redis-cache. Use a hosted Redis Cache (Azure, AWS, etc).
 - **Memcached.** Make sure you use pylibmc and not python-memcached, better performance with the C library.
- SSL still matters.

Azure Redis Cache

HOME

NOTIFICATIONS

BROWSE

ACTIVE

BILLING

New Redis Cache

DNS name
djangoapp ✓
.redis.cache.windows.net

PRICING TIER
Standard: 1 GB >

RESOURCE GROUP ⓘ
chatroom >

SUBSCRIPTION
Visual Studio Ultimate with MSDN >

LOCATION
North Central US >

Recommended pricing tiers

S1
STANDARD
102.67

B1
BASIC
40.92 ✓

S4
STANDARD
390.60

1 GB Cache

Dedicated service

SSL

Choose your pricing tier
BROWSE THE AVAILABLE PLANS AND THEIR FEATURES

<div>S0 Standard ✓</div> <div>250 MB Cache</div> <div>1 Replica</div> <div>Limited throughput</div> <div>Shared infrastructure</div> <div>SSL</div> <div>Replication</div> <div>99.9% SLA</div> <div>40.92 USD/MONTH (ESTIMATED)</div>	<div>S1 Standard</div> <div>1 GB Cache</div> <div>1 Replica</div> <div>Dedicated service</div> <div>SSL</div> <div>Replication</div> <div>99.9% SLA</div> <div>102.67 USD/MONTH (ESTIMATED)</div>	<div>S2 Standard</div> <div>2.5 GB Cache</div> <div>1 Replica</div> <div>Dedicated service</div> <div>SSL</div> <div>Replication</div> <div>99.9% SLA</div> <div></div>
<div>S3 Standard</div> <div>6 GB Cache</div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div>	<div>S4 Standard</div> <div>13 GB Cache</div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div>	<div>S5 Standard</div> <div>26 GB Cache</div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div>

Logging Django

Logging in Production

- Logstash
 - Collect logs, parse them, and store them for later use.
 - Free and open source.
- Set it up with Redis.

Logging in Production – Use Python-Logstash

```
LOGGING = {  
    'handlers': {  
        'logstash': {  
            'level': 'DEBUG',  
            'class': 'logstash.LogstashHandler',  
            'host': 'localhost',  
            'port': 5959, # Default value: 5959  
            'version': 1, # Version of logstash event schema. Default value: 0 (for backward  
compatibility of the library)  
            'message_type': 'logstash', # 'type' field in logstash message. Default value: 'logstash'.  
            'tags': ['tag1', 'tag2'], # list of tags. Default: None.  
        },  
    },  
},
```

Logging in Production – Use Python-Logstash

```
'loggers': {  
    'django.request': {  
        'handlers': ['logstash'],  
        'level': 'DEBUG',  
        'propagate': True,  
    },  
},  
}
```

Logging in Production – Using Sentry

- Realtime event logging and aggregation platform.
- Specialize in monitoring exceptions and errors.
- <https://github.com/getsentry/sentry>

PostgreSQL in Production

PostgreSQL in Production – pgpool & slony

- PostgreSQL 9+ has streaming replication assuming you're running identical databases with identical version on identical architectures (not really an issue).
- If not using the Streaming Replication feature (WAL), you'll want to use either pgpool or slony to scale your PostgreSQL database across multiple machines.
 - Use Slony if you want master-slave-peer replication.
 - Use pgpool or pgbouncer for connection pooling. Pg bouncer only does connection pooling.

Async Tasks & Django

Async Tasks – Celery

- Use Celery for Asynchronous Tasks
- Use Celery with Redis and store results in Django ORM
- RQ is an alternative for distributing tasks. (Based on Redis)

Read:

<http://www.caktusgroup.com/blog/2014/09/29/celery-production/>

On Redis

- Redis is extremely popular as an easy and simple publish-subscribe.
- Redis is very particular about memory. Running out of memory in Redis -> **CRASH AND BURN!**
- If you don't have enough memory in your VM to run Redis Cache -> Change your approach, use Apache Kafka.

Tips & Techniques

- Don't forget:
 - Turn off Debug Mode
 - Turn off Template Debug Mode
- Use different Django Settings Modules depending on the role and environment you are in through environment variables by setting "DJANGO_SETTINGS_MODULE".

What did we learn?


- Scaling Django Architectures
- Most of the work surrounds Django
- Minimal changes to your django app for the most part.

Thank You! Questions?

Follow @ramisayar

Resources

- <http://blog.disqus.com/post/62187806135/scaling-django-to-8-billion-page-views>
- <http://engineering.hackerearth.com/2013/10/07/scaling-database-with-django-and-haproxy/>
- <http://www.djangobook.com/en/2.0/chapter12.html>
- <https://www.digitalocean.com/community/tutorials/how-to-implement-ssl-termination-with-haproxy-on-ubuntu-14-04>
- <https://docs.djangoproject.com/en/1.7/topics/security/>
- <https://www.digitalocean.com/community/tutorials/how-to-scale-django-beyond-the-basics>



Microsoft

©2013 Microsoft Corporation. All rights reserved. Microsoft, Windows, Office, Azure, System Center, Dynamics and other product names are or may be registered trademarks and/or trademarks in the U.S. and/or other countries. The information herein is for informational purposes only and represents the current view of Microsoft Corporation as of the date of this presentation. Because Microsoft must respond to changing market conditions, it should not be interpreted to be a commitment on the part of Microsoft, and Microsoft cannot guarantee the accuracy of any information provided after the date of this presentation. MICROSOFT MAKES NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, AS TO THE INFORMATION IN THIS PRESENTATION.